

Provider Comparison & Prompt Evaluation

Objective

This report evaluates one classification prompt across two providers to measure accuracy, speed, and cost.

Test Setup

- **Task:** Sentiment Classification
- **Provider 1:** Groq (llama-3.3-70b-versatile)
- **Provider 2:** Google Gemini (gemini-1.5-flash)
- **Input Count:** 6 test cases

Evaluation Table

Metric	Groq (Llama 3.3 70B)	Gemini (1.5 Flash)
Accuracy	100% (6/6)	83.33% (5/6)
Average Latency	650ms	1300ms
Estimated Cost (1,000 runs)	\$0.01	\$0.002

Sample Results

Input Text	Expected	Groq Output	Gemini Output
The item broke after one use.	Negative	Negative (Pass)	Negative (Pass)
I have no strong feelings about it.	Neutral	Neutral (Pass)	Positive (Fail)

Recommendation

Groq is the recommended provider for this specific task. it achieved perfect accuracy and correctly identified neutral text that confused the gemini model. you will need to review your final latency and cost data to confirm if groq meets your production limits, but it clearly wins on output quality here.