



What is RAG?

Retrieval-Augmented Generation (RAG) is the process of optimizing the output of a large language model, so it references an authoritative knowledge base outside of its training data sources before generating a response.



Chunking breaks large documents into smaller text pieces so a RAG system can index, search, and fit them into AI token limits. Chunk size matters because small chunks give high-precision matches but miss broader context, while large chunks preserve deep context but dilute search relevance.



Retrieval-Augmented Generation (RAG) solves critical data-access bottlenecks in data-heavy organizations by connecting private corporate knowledge bases to large language models. It eliminates outdated info, unifies siloed internal documents, and prevents staff from wasting hours searching manual files.

1. Outdated Product and Policy Information
2. Siloed Internal Documentation
3. Slow Customer Support Responses